

SHORT TECHNICAL NOTE

New and simplified manual controls for projection and slice tours, with application to exploring classification boundaries in high dimensions

Ursula Laa^a, Alex Aumann^b, Dianne Cook^c, German Valencia^b

^aInstitute of Statistics, University of Natural Resources and Life Sciences, Vienna; ^bSchool of Physics and Astronomy, Monash University; ^cDepartment of Econometrics and Business Statistics, Monash University

ARTICLE HISTORY

Compiled September 23, 2022

ABSTRACT

This paper describes new user controls for examining high-dimensional data using low-dimensional linear projections and slices. A user can interactively change the contribution of a given variable to a low-dimensional projection, which is useful for exploring the sensitivity of structure to particular variables. The user can also interactively shift the center of a slice, for example, to explore how structure changes in local subspaces. Two Mathematica notebooks are provided, which contains functions enabling the user to experiment with these new manual controls, with one specifically for exploring regions and boundaries produced by classification models. The advantage of Mathematica is its linear algebra capabilities, and interactive cursor location controls. Some limited implementation has also been made available in the R package `tourr`.

KEYWORDS

data visualisation; grand tour; statistical computing; statistical graphics; multivariate data; dynamic graphics

1. Introduction

From a statistical perspective 3D is a rare data dimension, so unlike in most 3D rotation computer graphics applications, the more useful methods for data analysis need to work for arbitrary dimension. A good approach is to show projections from an arbitrary dimensional space to create dynamic data visualizations called *tours*. Tours involve views of high-dimensional (p) data with low-dimensional (d) projections. In his original paper on the grand tour, Asimov (1985) provided several algorithms for tour paths that could theoretically show the viewer the data *from all sides*. Prior to Asimov's work, there were numerous preparatory developments including Fisher-Keller, Friedman, and Tukey (1974)'s PRIM-9. PRIM-9 had user-controlled rotations on coordinate axes, allowing one to manually tour through low-dimensional projections. (A video illustrating the capabilities is available through video library of ASA Statistical Graphics Section (2022).) Steering through all possible projections is impossible, un-

CONTACT Ursula Laa. Email: ursula.laa@boku.ac.at, Alex Aumann. Email: aaum0002@student.monash.edu, Dianne Cook. Email: dicook@monash.edu, German Valencia. Email: german.valencia@monash.edu

like Asimov’s tours which allows one to quickly see many, many different projections. After Asimov there have been many tour developments, which are summarized in Lee et al. (2021).

One such direction of work develops the ideas from PRIM-9, to provide manual control of a tour. Cook and Buja (1997) describe controls for 1D (or 2D) projections, respectively in a 2D (or 3D) manipulation space, allowing the user to select any variable axis, and rotate it into, or out of, or around the projection through horizontal, vertical, oblique, radial or angular changes in value. Spyrisson and Cook (2020) refined this algorithm and implemented it to generate radial tour animation sequences.

Manual controls are especially useful for assessing sensitivity of structure to particular elements of the projection. There are many places where it is useful. In exploratory data analysis, where one sees clusters in a projection, one may ask whether some variables can be removed from the projection without affecting the clustering. For interpreting models, one can reduce or increase a variable’s contribution to examine the variable importance. Having the user interact with a projection is extremely valuable for understanding high-dimensional data. However, these algorithms have two problems: (1) the pre-processing of creating a manipulation space overly complicates the algorithm, (2) extending to higher dimensional control is difficult.

Another potentially useful manual control, is to allow the user to choose the position of the center of a slice. The slice tour was introduced in Laa, Cook, and Valencia (2020). It operates by converting the projection plane into a slice, by removing or de-emphasizing points that are further than a fixed orthogonal distance from the plane. The projection plane is usually thought of as passing through the center of the data. Manual control would allow the user to change the value of center point, by shifting it along a coordinate axis, while keeping the orientation of the projection plane fixed. The purpose would be to explore how or if the shape of the data, in the space orthogonal to the projection, changes as one gets away from the center. It would also allow the user to interactively decide on the thickness of the slice.

This paper explains the new manual controls for projection and slice tours. The next section describes the new algorithm for manual control, for both projections and slices. The use of these methods is illustrated to compare and contrast boundaries constructed by different classifiers. The software section describes a mathematica notebook that is used for the application, and describes the interactive environment that would be desirable within R as new technology becomes available. The paper is accompanied by an appendix with more details and adjustments to the manual controls, and two Mathematica notebooks, one focused on the application, and another that could be used generically for new data.

2. How to construct a manual tour

A manual tour allows the user to alter the coefficients of one (or more) variables contributing to a d dimensional projection. The initial ingredients are an orthonormal basis ($A_{p \times d}$) defining the projection of the data, and a variable id ($m \in \{1, \dots, p\}$) specifying which coefficient will be changed. A method to update the values of the component (m^{th} row of $A_{p \times d}$) of the controlled variable V_m is then needed.

2.1. Existing methods

The methods for updating component values in Cook and Buja (1997) (and utilised in Spyrison and Cook (2020)) are prescribed primarily for a 2D projection, to take advantage of (then) newly developed 3D trackball controls made available for computer gaming. The first step was to construct a 3D manipulation space from a 2D projection. In this space, the coefficient of the controlled variable ranges between -1 and 1. Movements of a cursor are recorded and converted into changes in the values of V_m thus changing the displayed 2D projection. The algorithm also provided constraints to horizontal, vertical, radial or angular motions only. The construction of the manipulation space overly complicates the manual controls, especially when considering possible techniques that will apply to arbitrary d .

2.2. A new simpler and broadly applicable approach

The new approach emerged from experiments on the tour using the linear algebra capabilities, and relatively new interactive graphics interface, available in Wolfram Research, Inc. (n.d.). The components corresponding to V_m are directly controlled by cursor movement, which updates row m of A . The updated matrix is then orthonormalised.

2.2.1. Algorithm

1. Provide A , and m . (Note that m could also be automatically chosen as the component that is closest to the cursor position.)
2. Change values in row m , for example, if $d = 2$ gives

$$A^* = [\mathbf{a}_1^* \ \mathbf{a}_2^*] = \begin{bmatrix} a_{11} & a_{12} \\ \vdots & \vdots \\ a_{m1}^* & a_{m2}^* \\ \vdots & \vdots \\ a_{p1} & a_{p2} \end{bmatrix}.$$

A large change in these values would correspond to making a large jump from the current projection. Small changes would correspond to tracking a cursor, making small jumps from the current projection.

3. Orthonormalise A^* , using Gram-Schmidt.

- i. Normalise \mathbf{a}_1^* and \mathbf{a}_2^* .
- ii. $\mathbf{a}_2^* = \mathbf{a}_2^* - (\mathbf{a}_1^* \cdot \mathbf{a}_2^*)\mathbf{a}_1^*$.

This algorithm will produce the changes to a projection as illustrated in Figure 1. The controlled variable, V_m , corresponds to the black line, and sequential changes to row m of A can be seen to roughly follow a specified position (orange dot). Changes in the other components happen as a result of the orthonormalisation, but are uncontrolled.




Figure 1. Sequence of projections where contribution of one variable is controlled (black) is changed using unconstrained orthonormalisation. The dot (orange) indicates the chosen values for the controlled variable. It can be seen that the actual axis does not precisely match the chosen position, but it is close.

2.3. Refinements to enforce exact position

The problem with the new simple method is that it is not faithful to the precise values for V_m because the orthonormalisation will change them. Even though these changes are for the most part imperceptible, one may wish to avoid them and there are numerous ways that this can be enforced, a few are detailed in the Appendix. These primarily differ in how the remaining variables are adjusted during orthonormalisation.

2.4. Manual control for slices

To better explore the space we combine the manual controls for the projection with manual controls for slicing. A slice is a section of the data that is defined by a projection, a center point that is anchoring it in the high-dimensional space and the slice thickness h (Laa, Cook, and Valencia 2020). A data point is inside the slice if its orthogonal distance from the projection plane (passing through the center point) is below the thickness h . This orthogonal distance is computed in terms of the component that is normal on the projection plane. For \mathbf{x}_i a p dimensional data point and \mathbf{c} the center point (in the same p dimensional space) we compute the orthogonal distance as

$$v_i^2 = \|\mathbf{x}'_i - \mathbf{c}'\|^2 = \mathbf{x}'_i{}^2 + \mathbf{c}'^2 - 2\mathbf{x}'_i \cdot \mathbf{c}', \quad (1)$$

with $\mathbf{c}' = \mathbf{c} - (\mathbf{c} \cdot \mathbf{a}_1)\mathbf{a}_1 - (\mathbf{c} \cdot \mathbf{a}_2)\mathbf{a}_2$, $\mathbf{x}'_i = \mathbf{x}_i - (\mathbf{x}_i \cdot \mathbf{a}_1)\mathbf{a}_1 - (\mathbf{x}_i \cdot \mathbf{a}_2)\mathbf{a}_2$ and $\mathbf{a}_k, k = 1, 2 (= d)$ denoting the columns of the projection matrix, $\mathbf{A} = (\mathbf{a}_1, \mathbf{a}_2)$.

2.4.1. Shifting the center

A natural starting point is to place \mathbf{c} in the center of the data distribution, but shifting it away from the mean can provide additional insights. In the case of a single orthogonal direction on the projection plane we can pick a sequence of center points \mathbf{c} in steps along that direction to move the slice and fully cover the data space. This no longer works in higher-dimensional spaces, and we can think of picking one direction and shifting the slice along the component orthogonal to the projection plane.

2.4.2. Changing the thickness

In addition it is also useful to interactively change the slice thickness h (also called the slice radius), in particular to find the preferred value for exploring the input data. For guidance the estimates of the number of points inside the slice as a function of the original sample size N and the number of dimensions p from Laa et al. (2022)

can be used: in case of a uniform distribution inside a sphere of radius R a slice with thickness h will contain N_S points, with

$$N_S(h, p, R, N) = \frac{N}{2} \left(\frac{h}{R} \right)^{p-2} \left(p - (p-2) \left(\frac{h}{R} \right)^2 \right). \quad (2)$$

3. Software

The implementation of the manual tour as suggested here requires the visualization of the current projection in terms of an axis display (see Figure 1 for an example). This display should be able to track the mouse position and adjust the projection based on the user selection. The user can select one axis (corresponding to one variable) by clicking on it, and then adjust the position of that axis by dragging it during the click event. In practice the closest axis will be selected, and the dragging results in tracking of the mouse position which will update the projection in small steps, such that no further interpolation is required.

The interactions in the axis display need to be mapped back onto the current projection matrix, which will then be orthonormalized before feeding back into the axis display. A second, linked display shows the projected data in sync with the updates from the axis display.

In addition, we also want to be able to look at slices of the data and select slicing parameters interactively. Here we have implemented a switch to change between projection and slice view, a slider that adjusts the slice height h and a numeric input to specify the center point c explicitly. The display will update based on these inputs, in particular the view will jump to a new slice when c is changed (though an interpolation might also be useful).

One might implement such an interactive interface via R Shiny, but in particular the tracking of the mouse position in small increments might pose a challenge. Here we found the dynamic graphics interface from Mathematica to be especially useful. In the following we first describe the relevant Mathematica functionalities and how they were used in the implementation, followed by a brief sketch of how the new approach might be implemented in the `tourr` package.

3.1. *Mathematica notebook*

Mathematica provides much utility and versatility, such as inbuilt data visualisation, data manipulation and analysis, dynamic functionality, and symbolic and numeric computation. Most of this inbuilt functionality is user-friendly and described in the Mathematica documentation; typically, numerous examples are provided within the documentation, and some possible issues are outlined there.

Importantly for our work, it is relatively simple to create dynamic objects via the inbuilt commands `Manipulate` or `DynamicModule`. Control objects, such as sliders, locator panes, and input fields, can then be used on dynamic variables, and when there is a change in the dynamic variable the dynamic objects which contain that variable will be updated. This is the essential ingredient for this implementation of the manual tour.

The most relevant inbuilt Mathematica function for our purpose is the locator pane. This creates a region on the screen where the position of the mouse is captured and

then converted to input that updates the graphics functions, enabling the manual navigation of the tour.

The examples presented in the application section below (and attached in notebook format in the supplementary material) all use our primary new function, `SliceDynamic`. This function typically accepts grouped data in the form of a matrix where the second last column details the name of the group and the last column details the group index (which can be one if there is only one group). After specifying the initial slice thickness and the slice range, the user is presented with an interactive display in which the control objects appear on the left and the slice visualisation appears on the right. The user can change the orientation of the slice via the locator pane, which changes the projection matrix; a slider controls the slice thickness; and there is an input field, which changes the slice centre. The user can also change the appearance of the plot by zooming into the center or changing the point size with sliders provided. It is worth noting that this zoom works best with data that has been centred and scaled. The projected data can be displayed to contrast it with the slice by ticking a box. The explicit projection matrix can also be displayed via another checkbox. This is especially useful when we wish to import a projection that was identified as interesting in the manual exploration later, as shown in the applications below.

Details about the implementation and usage instructions are given in the Appendix.

3.2. *Extensions to the R package tourr*

The R package `tourr` (Wickham et al. 2011) provides numerous types of tours. An additional function, `radial_tour()` has been constructed that will generate a sequence of projections that will decrease the coordinates for V_m to $\mathbf{0}$ and back to the original values. It is applicable for any projection dimension, d .

Changing slice center

3.3. *Desirable interactivity*

Potential adjustments for future work include a better visualization of the current center point, potentially with interactive graphics to move the slice similar to how the manual tour was implemented here. An alternative would be a step-wise shift of the slice towards the new center point to get a better understanding of the distribution along the selected direction.

4. Application

To illustrate the usefulness of the manual controls we use the 4D penguins data (Horst, Hill, and Gorman 2020) and look at classification models following Wickham, Cook, and Hofmann (2015). We will show how classification boundaries can be explored and better understood on projections and slices through 4D space. Figure 2 shows a scatterplot matrix of this data. There are four variables (`b1 = bill_length_mm`, `bd = bill_depth_mm`, `f1 = flipper_length_mm`, `bm = body_mass_g`) measuring the size of the penguins from three species (Adelie, Chinstrap and Gentoo). The scatterplot matrix shows that the three species appear to be likely separable, and that at least the Gentoo can be distinguished from the other two species when `bd` is paired with `f1` or `bm`. The steps for exploring boundaries in this example are as follows:

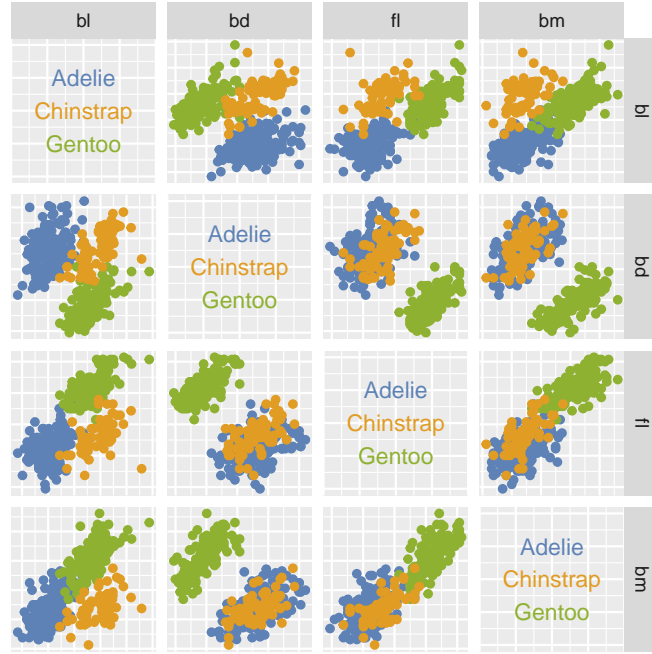


Figure 2. Scatterplot matrix of the (standardised) penguins data. The three species are reasonably different in size, with Gentoo distinguished from the other two on body depth relative to flipper length and body mass.

1. Build your classification model.
2. Predict the class for a dense grid of values covering the data space.
3. Examine projections, using a manual tour so that the contribution of any variable is controlled.
4. Slice through the center, to explore where the boundaries will likely meet.
5. Move the slice by changing the center in the direction of a single variable to explore the extent of a boundary for a single group relative to a variable.

4.1. *Constructing the 4D prediction regions*

We use the `classifly` package (Wickham 2022) to generate predictions across the 4D cube spanned by the data, with two classification models: linear discriminant analysis (LDA) and random forest (RF). Both the data and the model points in the grid are centered and scaled (standard deviation = 1).

4.2. *Exploring projections manually*

We start by first exploring the projections of the model prediction. Figure 3 summarises the process. Through manual rotation of the view we can get a feeling for where in the space we primarily predict each of the three species, and we also get a sense of the difference between the two models. To illustrate this difference we have manually rotated the projection for the RF model (left plot) to identify a view that shows the non-linear but block-type structure that is typical of this type of model. This particular projection (A_1) is exported so that the same projection can be used to show the LDA model (middle plot) and the actual data (right plot). What can be seen is the linearity of the LDA model, where the boundaries are linear and oblique to the variable axes. And, interestingly this particular projection of the original data shows very distinct

clusters of the three species. That means, the obscuring of the boundaries between groups for both of the models is driven by what is happening in the orthogonal space to the plane of the selected projection.

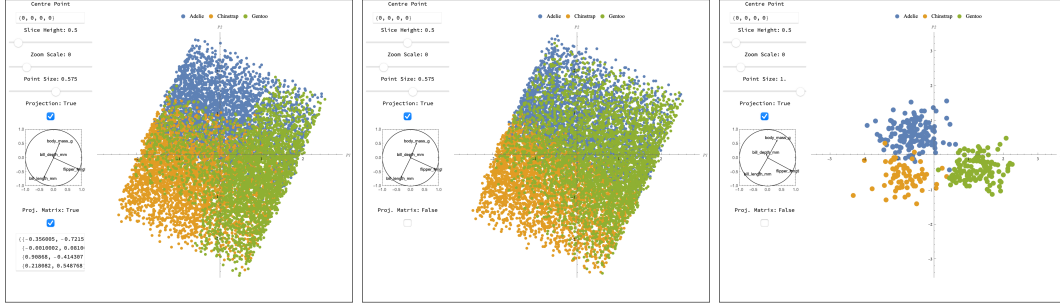


Figure 3. Projection identified using the manual tour, because it reveals an interesting structure in the predictions from the RF model (left). We can clearly see a block structure, while the LDA model (middle) produces linear boundaries. The three groups are nicely separated in this projection of the data (right).

This example can be reproduced with the code in `penguins_exploring_manually.nb` and a run through is shown in the video at <https://vimeo.com/747585410>.

4.3. *Slicing through the center*

We continue the investigation by now slicing orthogonally to the projection A_1 . For both models we look at a thin ($h = 0.5$) slice through the center, $\mathbf{c}^0 = (0, 0, 0, 0)$. At first we explore how changing the projection away from A_1 can help with understanding the boundary better. For our example, notice that A_1 does not contain any contribution from the second variable (`bd`), so we will first rotate this variable into the view.

Figure 4 shows snapshots of the exploration. The top row is the initial projection (A_1 and \mathbf{c}^0), and the bottom row is a later projection containing more of `bd` (new projection A_2 while keeping the center fixed at \mathbf{c}^0). The columns correspond to the two models, with RF on the left and LDA on the right.

Both models have some overlap at the center for the thin slice based on A_1 (despite the small value for h). The second slice (based on A_2) mostly resolves this overlap and reveals the primary differences between the models. The boundary between Adelie (blue) and Chinstrap (yellow) is very similar for both models. The boundary between Chinstrap and Gentoo (green) is where they differ.

The RF is almost straight in this view, as something we might expect from a tree model where splitting occurs on single variables only. However, it is straight in a combination of the second and third variable (`bd` and `f1`). By examining the scatterplot matrix, we can understand how this boundary was built: on each of `bd` and `f1` it is possible to cut on a value where most of the Gentoo penguins are different from the other two species for any tree, and likely only one is needed. Thus the forest construction is providing a sample of trees that use one variable or the other variable to split, producing the blocky boundary on the combination of the two.

In contrast, the LDA boundary has an oblique split between the two, and roughly divides the space into three similarly sized areas. It is almost like a textbook illustration of how LDA works for two variables (2D) where by assuming the clusters have equal variance-covariance, it places a boundary half-way between the group means.

Finally, to determine which model does the boundaries better, compare them with the A_2 projection of the data (Figure 5). Gentoo is more separated from the other



Figure 4. Comparing slices based on two projections A_1 (top row) and A_2 (bottom row), for the two models RF (left) and LDA (right). With A_1 we see two groups overlap (green - Gentoo with yellow - Chinstrap), while the rotation to A_2 results in clear boundaries inside the slice. The boundary between Adelie (blue) and Chinstrap (yellow) is similar for both models but very different between Chinstrap and Gentoo (green).

two in this projection, and one can imagine that the trees in the forest has greedily grasped any one of many places to make a split to separate the group. It might be argued though that the RF boundary is cut too close to the Chinstrap species, and might lead to some unnecessary misclassification with new data. The LDA boundary is better placed for all species.

This example can be reproduced with the code in `pen-guins_slicing_through_center.nb` and a run through is shown in the first part of the video at <https://vimeo.com/747590472>.

4.4. *Shifting the slice center*

We have seen that starting from A_1 using the manual controls to change the contribution of the second variable we could find a clear separation boundary indicating the relation between this variable (bill depth) and the Gentoo penguin species. Instead of rotating to a different projection, we might also change the view by moving the slice along one axis in the 4D space. Here we will continue our exploration of the dependence on `bd` and move the slice defined by A_1 to either large positive or negative values ($c_{bd}^{\pm} = \pm 1.5$ after centering and scaling). Thus we define $\mathbf{c}^{\pm} = (0, c_{bd}^{\pm}, 0, 0)$. Here we will also look at slices of the observed data points, using a thicker slice ($h = 1.5$) to




Figure 5. Projection of the data based on A_2 . Compared to projecting onto A_1 we see that the green observations (Gentoo) are more separated from the other two species.

capture enough points in a given view.

We start by a comparison of the two models and the data distribution shifting to \mathbf{c}^+ , thus the slice is localized towards high values of bill depth in Figure 6. We can see that all three slices (the two models and the data) contain almost no points from the third class (green, Gentoo), and that the decision boundary between the two models is very similar.




Figure 6. Shifting the slice center in the positive direction of bill depth produces regions that have no Gentoo (green). The two models have similar boundaries except that the RF is more of a step function.

A more interesting comparison is found for \mathbf{c}^- , thus the slice localized towards low values of bill depth, shown in Figure 7. The RF model (left) predicts all three species within this slice, with an interesting boundary for the third class (green, Gentoo). On the other hand the LDA model (middle) predominantly predicts the third class within the slice, this appears to be enforced through the linear structure of the model. Looking finally at the thick slice through the data we see that there are primarily observations from this class within the slice we can conclude that the two models have filled in the “empty” space (where we do not have any training observations) in very different ways and according to what we might expect given the model structure.

Finally it is interesting to compare the slice views to the projection of the models seen in Fig. 3 to better understand how the boundaries change along the bd direction and where the differences in the projections come from.

This example can be reproduced with the code in `penguins_shifting_slice_center.nb` and a run through is shown in the second part of the video at <https://vimeo.com/747590472>.




Figure 7. Shifting the slice center in the negative direction of bill depth produces regions that are mostly Gentoo (green). The two models have very different boundaries: the nonlinear potential of RF can be seen here where still some subspaces would predict to be Adelie and Chinstrap.

5. Discussion

This short note has described new technology for manually interacting with tours. The manual control of projection coefficients is most useful for assessing variable importance to perceived structure, but can be generally used for steering a viewer through high-dimensional space. Changing the center of a slice manually enables exploring the space orthogonal to a projection, specifically in the direction of a single variable. This is a new tool, and as shown in the example, when used with the slice tour can be very useful for understanding boundaries of classifiers. This is also likely a useful method for examining low-dimensional non-linear manifolds, and also functions of multiple parameters.

Mathematica provided a useful sandbox to experiment with the ideas presented in this paper. Most of the data, though, is first constructed using R, especially the classification boundaries. If new technology for interactive graphics becomes available in R it would be useful to create a tighter coupling of models and visualization to allow exploring and comparing fits.

Acknowledgements

The authors gratefully acknowledge the support of the Australian Research Council and the ResearchFirst undergraduate research program at Monash University. The paper was written in `rmarkdown` (Xie, Allaire, and Golemund 2018) using `knitr` (Xie 2015). The scatterplot matrix of the penguins data was produced by the `GGally` (Emerson et al. 2013) package built on `ggplot2` (Wickham 2016) graphics. We thank the Institute of Statistics, BOKU, for their hospitality while part of this work was conducted.

Supplementary material

The source material and animated gifs for this paper are available at
 Mathematica notebook
 Appendix

References

- ASA Statistical Graphics Section. 2022. “Video Library.” <https://community.amstat.org/jointscsg-section/media/videos>.
- Asimov, D. 1985. “The Grand Tour: A Tool for Viewing Multidimensional Data.” *SIAM Journal of Scientific and Statistical Computing* 6 (1): 128–143.
- Cook, Dianne, and Andreas Buja. 1997. “Manual Controls for High-Dimensional Data Projections.” *Journal of Computational and Graphical Statistics* 6 (4): 464–480. <http://www.jstor.org/stable/1390747>.
- Emerson, John W., Walton A. Green, Barret Schloerke, Jason Crowley, Dianne Cook, Heike Hofmann, and Hadley Wickham. 2013. “The Generalized Pairs Plot.” *Journal of Computational and Graphical Statistics* 22 (1): 79–91. Accessed 2022-07-14. <http://www.jstor.org/stable/43304816>.
- FisherKeller, M.A., J.H. Friedman, and J.W. Tukey. 1974. “PRIM-9, an interactive multidimensional data display and analysis system.” In *The Collected Works of John W. Tukey: Graphics 1965-1985, Volume V*, edited by William S. Cleveland, 340–346.
- Horst, Allison Marie, Alison Presmanes Hill, and Kristen B Gorman. 2020. *palmerpenguins: Palmer Archipelago (Antarctica) penguin data*. R package version 0.1.0, <https://allisonhorst.github.io/palmerpenguins/>.
- Laa, Ursula, Dianne Cook, Andreas Buja, and German Valencia. 2022. “Hole or Grain? A Section Pursuit Index for Finding Hidden Structure in Multiple Dimensions.” *Journal of Computational and Graphical Statistics* 0 (0): 1–14. <https://doi.org/10.1080/10618600.2022.2035230>.
- Laa, Ursula, Dianne Cook, and German Valencia. 2020. “A Slice Tour for Finding Hollowness in High-Dimensional Data.” *Journal of Computational and Graphical Statistics* 29 (3): 681–687. <https://doi.org/10.1080/10618600.2020.1777140>.
- Lee, Stuart, Dianne Cook, Natalia da Silva, Ursula Laa, Earo Wang, Nick Spyrisson, and H. Sherry Zhang. 2021. “Advanced Review: The State-of-the-Art on Tours for Dynamic Visualization of High-dimensional Data.” *arXiv:2104.08016 [cs, stat]* <http://arxiv.org/abs/2104.08016>.
- Spyrisson, Nicholas, and Dianne Cook. 2020. “spinifex: an R Package for Creating a Manual Tour of Low-dimensional Projections of Multivariate Data.” *The R Journal* 12 (1): 243. <https://journal.r-project.org/archive/2020/RJ-2020-027/index.html>.
- Wickham, Hadley. 2016. *ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York. <https://ggplot2.tidyverse.org>.
- Wickham, Hadley. 2022. *classify: Explore Classification Models in High Dimensions*. R package version 0.4.1, <https://CRAN.R-project.org/package=classify>.
- Wickham, Hadley, Dianne Cook, and Heike Hofmann. 2015. “Visualizing statistical models: Removing the blindfold.” *Statistical Analysis and Data Mining: The ASA Data Science Journal* 8 (4): 203–225. <https://onlinelibrary.wiley.com/doi/abs/10.1002/sam.11271>.
- Wickham, Hadley, Dianne Cook, Heike Hofmann, and Andreas Buja. 2011. “tourr: An R Package for Exploring Multivariate Data with Projections.” *Journal of Statistical Software* 40 (2): 1–18. <http://www.jstatsoft.org/v40/i02/>.
- Wolfram Research, Inc. n.d. “Mathematica, Version 13.1.” Champaign, IL, 2022, <https://www.wolfram.com/mathematica>.
- Xie, Yihui. 2015. *Dynamic Documents with R and knitr*. 2nd ed. Boca Raton, Florida: Chapman and Hall/CRC. <https://yihui.name/knitr/>.
- Xie, Yihui, Joseph J. Allaire, and Garrett Golemund. 2018. *R Markdown: The Definitive Guide*. Chapman and Hall/CRC. <https://bookdown.org/yihui/rmarkdown>.